

□ Tuned LSTM Stock Movement Predictor

This notebook loads stock price data, creates technical indicators, prepares the data for LSTM, builds a tuned LSTM model, and evaluates its performance.

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
from sklearn.preprocessing import MinMaxScaler
from sklearn.metrics import accuracy_score, classification_report
from sklearn.model_selection import train_test_split
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import LSTM, Dense, Dropout
from tensorflow.keras.callbacks import EarlyStopping
from tensorflow.keras.optimizers import Adam
import ta # Technical Analysis library

# Load dataset
df = pd.read_csv("a.csv")
df['Date'] = pd.to_datetime(df['Date'])
df.sort_values('Date', inplace=True)

# Add technical indicators
df['rsi'] = ta.momentum.RSIIndicator(df['Close']).rsi()
df['macd'] = ta.trend.MACD(df['Close']).macd()
df['ema'] = ta.trend.EMAIndicator(df['Close']).ema_indicator()

# Create features & target
df['target'] = np.where(df['Close'].shift(-1) > df['Close'], 1, 0)
df.dropna(inplace=True)

features = ['Open', 'High', 'Low', 'Close', 'Volume', 'rsi', 'macd', 'ema']
data = df[features]

# Scale features
scaler = MinMaxScaler()
scaled_data = scaler.fit_transform(data)

# Create sequences
sequence_length = 60
X = []
y = []
for i in range(sequence_length, len(scaled_data)):
    X.append(scaled_data[i-sequence_length:i])
    y.append(df['target'].values[i])
X = np.array(X)
y = np.array(y)
```

```

# Train-test split
X_train, X_test, y_train, y_test = train_test_split(
    X, y, test_size=0.1, shuffle=False
)

# LSTM Model
model = Sequential()
model.add(LSTM(128, return_sequences=True,
input_shape=(X_train.shape[1], X_train.shape[2])))
model.add(Dropout(0.3))
model.add(LSTM(64))
model.add(Dropout(0.3))
model.add(Dense(1, activation='sigmoid'))

optimizer = Adam(learning_rate=0.0005)
model.compile(loss='binary_crossentropy', optimizer=optimizer,
metrics=['accuracy'])

early_stop = EarlyStopping(monitor='val_loss', patience=7,
restore_best_weights=True)

history = model.fit(
    X_train, y_train,
    epochs=100,
    batch_size=64,
    validation_data=(X_test, y_test),
    callbacks=[early_stop],
    verbose=1
)

```

Epoch 1/100

C:\Users\user\AppData\Roaming\Python\Python312\site-packages\keras\src\layers\rnn\rnn.py:200: UserWarning: Do not pass an `input_shape`/`input_dim` argument to a layer. When using Sequential models, prefer using an `Input(shape)` object as the first layer in the model instead.

```
super().__init__(**kwargs)
```

```
93/93 ————— 11s 99ms/step - accuracy: 0.5072 - loss:
0.6943 - val_accuracy: 0.5030 - val_loss: 0.6949
```

Epoch 2/100

```
93/93 ————— 8s 83ms/step - accuracy: 0.4982 - loss:
0.6940 - val_accuracy: 0.5030 - val_loss: 0.6949
```

Epoch 3/100

```
93/93 ————— 38s 407ms/step - accuracy: 0.5008 - loss:
0.6944 - val_accuracy: 0.5030 - val_loss: 0.6933
```

Epoch 4/100

```
93/93 ————— 118s 1s/step - accuracy: 0.5059 - loss:
0.6933 - val_accuracy: 0.5030 - val_loss: 0.6934
```

```

Epoch 5/100
93/93 _____ 131s 1s/step - accuracy: 0.5011 - loss:
0.6933 - val_accuracy: 0.5030 - val_loss: 0.6939
Epoch 6/100
93/93 _____ 51s 552ms/step - accuracy: 0.5117 - loss:
0.6927 - val_accuracy: 0.5030 - val_loss: 0.6931
Epoch 7/100
93/93 _____ 7s 80ms/step - accuracy: 0.4850 - loss:
0.6940 - val_accuracy: 0.4970 - val_loss: 0.6939
Epoch 8/100
93/93 _____ 8s 84ms/step - accuracy: 0.5130 - loss:
0.6933 - val_accuracy: 0.5030 - val_loss: 0.6933
Epoch 9/100
93/93 _____ 8s 82ms/step - accuracy: 0.5070 - loss:
0.6936 - val_accuracy: 0.4894 - val_loss: 0.6933
Epoch 10/100
93/93 _____ 7s 80ms/step - accuracy: 0.4979 - loss:
0.6938 - val_accuracy: 0.5030 - val_loss: 0.6933
Epoch 11/100
93/93 _____ 8s 83ms/step - accuracy: 0.5136 - loss:
0.6929 - val_accuracy: 0.4970 - val_loss: 0.6938
Epoch 12/100
93/93 _____ 8s 83ms/step - accuracy: 0.5145 - loss:
0.6927 - val_accuracy: 0.5030 - val_loss: 0.6932
Epoch 13/100
93/93 _____ 7s 77ms/step - accuracy: 0.5057 - loss:
0.6933 - val_accuracy: 0.5000 - val_loss: 0.6934

```

Evaluation

```

y_pred = (model.predict(X_test) > 0.5).astype("int32")
print("Accuracy:", accuracy_score(y_test, y_pred))
print("Classification Report:\n", classification_report(y_test,
y_pred))
# Plot accuracy
plt.plot(history.history['accuracy'], label='Train Accuracy')
plt.plot(history.history['val_accuracy'], label='Validation Accuracy')
plt.legend()
plt.title('Model Accuracy over Epochs')
plt.show()

```

21/21 _____ 1s 25ms/step


Accuracy: 0.5030211480362538

Classification Report:

	precision	recall	f1-score	support
0	0.00	0.00	0.00	329
1	0.50	1.00	0.67	333
accuracy			0.50	662
macro avg	0.25	0.50	0.33	662

weighted avg 0.25 0.50 0.34 662

```
C:\Users\user\AppData\Roaming\Python\Python312\site-packages\sklearn\
metrics\_classification.py:1565: UndefinedMetricWarning: Precision is
ill-defined and being set to 0.0 in labels with no predicted samples.
Use `zero_division` parameter to control this behavior.
    _warn_prf(average, modifier, f"{metric.capitalize()} is",
len(result))
C:\Users\user\AppData\Roaming\Python\Python312\site-packages\sklearn\
metrics\_classification.py:1565: UndefinedMetricWarning: Precision is
ill-defined and being set to 0.0 in labels with no predicted samples.
Use `zero_division` parameter to control this behavior.
    _warn_prf(average, modifier, f"{metric.capitalize()} is",
len(result))
C:\Users\user\AppData\Roaming\Python\Python312\site-packages\sklearn\
metrics\_classification.py:1565: UndefinedMetricWarning: Precision is
ill-defined and being set to 0.0 in labels with no predicted samples.
Use `zero_division` parameter to control this behavior.
    _warn_prf(average, modifier, f"{metric.capitalize()} is",
len(result))
```



```
def predict_tomorrow_movement(model, scaler):
    # Manually input today's values
```

```

open_price = 243.8
high_price = 249.2
low_price = 240.7
close_price = 244.3
volume = 1234567

# Technical indicators – dummy values or real ones you compute
rsi = 51.2
macd = 0.5
ema = 245.1

# Combine in the same order as used during training
input_features = np.array([[open_price, high_price, low_price,
close_price,
                                volume, rsi, macd, ema]])

# Scale the input
input_scaled = scaler.transform(input_features)

# Reshape for LSTM (3D): We have 1 sample, 60 timesteps, and 8
features.
# Since we have only one data point, we'll repeat it across the 60
timesteps.
input_scaled = np.repeat(input_scaled, 60, axis=0) # Repeat the
input for 60 timesteps
input_scaled = input_scaled.reshape((1, 60,
input_scaled.shape[1])) # Shape: (1, 60, 8)

# Predict
prediction = model.predict(input_scaled)
result = int(prediction[0] > 0.5)

# Output
print("\n Prediction for Tomorrow:")
print("→ Price expected to go UP " if result == 1 else "← Price
expected to go DOWN ")

# Call the prediction function after training
predict_tomorrow_movement(model, scaler)

```

1/1 ————— 0s 36ms/step

```

 Prediction for Tomorrow:
← Price expected to go DOWN

```

C:\Users\user\AppData\Roaming\Python\Python312\site-packages\sklearn\utils\validation.py:2739: UserWarning: X does not have valid feature names, but MinMaxScaler was fitted with feature names

```
warnings.warn(
C:\Users\user\AppData\Local\Temp\ipykernel_39616\1557592171.py:28:
```

DeprecationWarning: Conversion of an array with ndim > 0 to a scalar is deprecated, and will error in future. Ensure you extract a single element from your array before performing this operation. (Deprecated NumPy 1.25.)

```
result = int(prediction[0] > 0.5)
```